

# Measures of association Homework

Tim Hulak

I produced the material below with no assistance

## Exercises 3,4, 8, 9, and 10 on pages 155 and 156 of *Reasoning with Data: An Introduction to Traditional and Bayesian Statistics Using R*

3. Run `cor.test()` on the correlation between “area” and “perm” in the rock data set and interpret the results. Note that you will have to use the “`jaccessortogetateachofthetwovariables(likethis : rockarea)`”. Make sure that you interpret both the confidence interval and the p-value that is generated by `cor.test()`.

```
options(scipen=999)
data("rock")
head(rock)
```

```
##   area    peri    shape perm
## 1 4990 2791.90 0.0903296  6.3
## 2 7002 3892.60 0.1486220  6.3
## 3 7558 3930.66 0.1833120  6.3
## 4 7352 3869.32 0.1170630  6.3
## 5 7943 3948.54 0.1224170 17.1
## 6 7979 4010.15 0.1670450 17.1
```

```
cor.test(rock$area,rock$perm)
```

```
##
## Pearson's product-moment correlation
##
## data:  rock$area and rock$perm
## t = -2.9305, df = 46, p-value = 0.005254
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
##  -0.6118206 -0.1267915
## sample estimates:
##          cor
## -0.396637
```

**Answer:** The output for the *Pearson's product-moment correlation* has a *lower-bound confidence interval* of **-0.6118206** and an *upper-bound confidence interval* of **-0.1267915**. Since both of these values are negative, it can be safely assumed that the correlation is skewed in a particular direction. The *p-value* is **0.005254**, which is less than the traditional alpha of **0.05**. This means that we can reject the null hypothesis that the correlation between the rock area and the rock perm is not 0.

- 
4. Create a copy of the `bfCorTest()` custom function presented in this chapter. Don't forget to "source" it (meaning that you have to run the code that defines the function one time to make R aware of it). Conduct a Bayesian analysis of the correlation between "area" and "perm" in the rock data set.

```
#Copied from the book:  
# Reasoning with Data: An Introduction to Traditional and Bayesian Statistics Using R  
# Stanton, Jeffrey M.. Reasoning with Data (p. 136). Guilford Publications. Kindle Edition.  
library("BayesFactor")
```

```
## Loading required package: coda
```

```
## Loading required package: Matrix
```

```
## *****
```

```
## Welcome to BayesFactor 0.9.12-4.2. If you have questions, please contact Richard Morey (richarddmorey@ucsd.edu)
```

```
##
```

```
## Type BFManual() to open the manual.
```

```
## *****
```

```
bfCorTest <- function(x,y) #Get r from BayesFactor  
{  
  zx <- scale(x) #Standardize X  
  zy <- scale(y) #Standardize Y  
  zData <- data.frame(x=zx,rhoNot0=zy) #Put in a data frame  
  bfOut <- generalTestBF(x ~ rhoNot0, data=zData) #linear coefficient  
  mcmcOut <- posterior(bfOut,iterations=10000) #posterior samples  
  print(summary(mcmcOut[, "rhoNot0"])) #Show the HDI for r  
  return(bfOut) #Return Bayes factor object  
}
```

```
bfCorTest(rock$area,rock$perm)
```

```
##
```

```
## Iterations = 1:10000
```

```
## Thinning interval = 1
```

```
## Number of chains = 1
```

```
## Sample size per chain = 10000
```

```
##
```

```
## 1. Empirical mean and standard deviation for each variable,  
##    plus standard error of the mean:
```

```
##
```

```
##           Mean           SD      Naive SE Time-series SE  
##    -0.348678     0.136155    0.001362      0.001498
```

```
##
```

```
## 2. Quantiles for each variable:
```

```
##
```

```
##      2.5%      25%      50%      75%     97.5%  
## -0.62324 -0.43914 -0.34756 -0.25832 -0.08366
```

```
## Bayes factor analysis
```

```
## -----
## [1] rhoNot0 : 8.072781 ±0%
##
## Against denominator:
##   Intercept only
## ---
## Bayes factor type: BFlinearModel, JZS
```

**Answer:** The point estimate for rho is **is -0.3477**. The *95% Confidence Interval* has a lower-bound of **-0.61125** and an upper-bound of **-0.08239**. This range does not straddle zero. Finally, the Bayes Factor is approximately **8.07**. This means that there is a ratio of **0.1239157:1** odds in favor of the null hypothesis. In other words, there is strong evidence to reject the null hypothesis in favor if the alternative hypothesis (especially when the a rule-of-thumb odds cutoff value of 3:1 is considered).

- 
8. Not unexpectedly, there is a data set in R that contains these data. The data set is called UCBA admissions and you can access the department mentioned above like this: UCBAAdmissions[, ,1]. Make sure you put two commas before the 1: this is a three dimensional contingency table that we are subsetting down to two dimensions. Run `chisq.test()` on this subset of the data set and make sense of the results.

```
data('UCBAAdmissions')
UCBAAdmissions[, ,1]
```

```
##           Gender
## Admit      Male Female
##   Admitted  512     89
##   Rejected  313     19
```

```
chisq.test(UCBAAdmissions[, ,1], correct=F)
```

```
##
## Pearson's Chi-squared test
##
## data:  UCBAAdmissions[, , 1]
## X-squared = 17.248, df = 1, p-value = 0.0000328
```

**Answer:** The output of the *Pearson's Chi-squared test* shows a *chi-squared* value of **17.248** with only **1 degree of freedom**. The *p-value* is **0.0000328**, which is less than the traditional alpha of **0.05**. This means that we can reject the null hypothesis and conclude that gender and admit are not independent of each other.

- 
9. Use `contingencyTableBF()` to conduct a Bayes factor analysis on the UCB admissions data. Report and interpret the Bayes factor.

```
admission_table <- contingencyTableBF(UCBAAdmissions[, ,1], sampleType="poisson", posterior=F)
summary(admission_table)
```

```
## Bayes factor analysis
## -----
## [1] Non-indep. (a=1) : 1111.64 ±0%
##
## Against denominator:
##   Null, independence, a = 1
## ---
## Bayes factor type: BFcontingencyTable, poisson
```

**Answer:** There is a **0.0008996:1** ratio between the two factors. They are independent from one another (the two factors are not associated). Because Bayes factor exceeds the rule-of-thumb ratio of **3:1**, we can treat it as positive evidence in favor of independence.

- 
10. Using the UCBA data, run `contingencyTableBF()` with posterior sampling. Use the results to calculate a 95% HDI of the difference in proportions between the columns.

```
admission_table_posterior <- contingencyTableBF(UCBAAdmissions[, , 1], sampleType="poisson", posterior=T,
summary(admission_table_posterior)
```


```
##
## Iterations = 1:10000
## Thinning interval = 1
## Number of chains = 1
## Sample size per chain = 10000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean      SD Naive SE Time-series SE
## lambda[1,1] 511.40 22.743 0.22743      0.22743
## lambda[2,1] 312.79 17.505 0.17505      0.17505
## lambda[1,2]  89.68  9.580 0.09580      0.09445
## lambda[2,2]  19.96  4.486 0.04486      0.04282
##
## 2. Quantiles for each variable:
##
##           2.5%    25%    50%    75%   97.5%
## lambda[1,1] 467.98 495.74 511.26 526.47 557.07
## lambda[2,1] 279.19 300.62 312.44 324.47 347.59
## lambda[1,2]  72.01  83.08  89.38  96.05 109.35
## lambda[2,2]  12.20  16.73  19.60  22.80  29.69
```

```
head(admission_table_posterior)
```

```
## Markov Chain Monte Carlo (MCMC) output:
## Start = 1
## End = 7
## Thinning interval = 1
##      lambda[1,1] lambda[2,1] lambda[1,2] lambda[2,2]
## [1,]    525.0880    330.2580    83.31960    28.89413
```


```
## [2,] 498.9203 326.0525 81.86942 21.71752
## [3,] 516.9967 318.7159 90.17500 17.14817
## [4,] 529.0354 337.2086 88.04313 16.56355
## [5,] 471.2954 314.7650 105.86022 18.53270
## [6,] 510.0820 311.8600 81.69840 22.10735
## [7,] 499.2450 342.2604 88.19060 17.82879
```

```
male <- admission_table_posterior[, "lambda[1,1]"]/admission_table_posterior[, "lambda[2,1]"]
hist(male)
```



```
female <- admission_table_posterior[, "lambda[1,2]"]/admission_table_posterior[, "lambda[2,2]"]
hist(female)
```

## Histogram of female



```
difference <- male - female
head(difference)
```

```
## Markov Chain Monte Carlo (MCMC) output:
## Start = 1
## End = 7
## Thinning interval = 1
## [1] -1.293684 -2.239557 -3.636454 -3.746608 -4.214784 -2.059920 -3.487858
```

```
c(quantile(difference,c(0.025)),quantile(difference,c(0.975)))
```

```
##      2.5%      97.5%
## -5.986787 -1.213336
```

**Answer:** The 95% HDI of the difference in proportions between the columns appears to be between **-6.056048** and **-1.217753**.